

CS8656 Project Proposal // An Efficient Fill Estimation Algorithm for Sparse Tensors in Blocked Formats

Peter Ahrens, Nicholas Schiefer, Helen Xu

We present an algorithm to efficiently compute an important heuristic for autotuning sparse tensor operations. A tensor is a multidimensional array. A sparse tensor is a tensor whose entries are mostly zeros, which do not have to be stored or operated on in most linear algebraic operations. Since sparse tensors typically contain more than 90% zero entries, taking advantage of sparsity can provide substantial increases in performance. However, the increased complexity of datastructures that can describe the irregular locations of nonzeros in these tensors poses a significant challenge to performance engineers.

These challenges are magnified in an era of increasing heterogeneity of processors. In order to write the most efficient sparse tensor code, the programmer must take into account both the target architecture and the relevant structural properties of the nonzeros of the sparse tensor. Writing custom code for each processor requires extensive engineering effort and the structure of nonzeros is usually known only at runtime. Therefore, autotuning (automatically generating customized code) has become a necessary part of writing efficient sparse code.

Previous efforts in autotuning for sparse tensors focus on sparse matrices, which are more broadly applicable in fields ranging from scientific computing to machine learning. The diverse space of operations and nonzero patterns of sparse matrices have led to the development of a wide variety of sparse matrix formats that allow programmers to more efficiently operate on the matrices. We limit our description to perhaps the most popular such format, Compressed Sparse Row (CSR) and a variant we will call Blocked Compressed Sparse Row (BCSR). In CSR, only the nonzeros and their locations are stored in each row of the matrix. In Blocked Compressed Sparse Row, an $m \times n$ matrix is divided into $m/r \times n/c$ submatrices, where each submatrix is of size $r \times c$. Only blocks which contain nonzeros are stored, and these nonzero blocks are stored in CSR format. Storing the locations of nonzero blocks requires less memory and less computational logic than storing the locations of individual nonzeros. For matrices that naturally have a block structure of nonzeros, such as those produced by finite element methods, this can lead to an increase in the performance of sparse matrix operations. However, because we store the entire block (filling in zero entries with explicit zeros) setting the block size to be too large can decrease the performance.

Given the definition of BCSR, it is natural to wonder how one might choose the correct block size for a given matrix. Vuduc et. al. describes an effective heuristic for predicting the performance P (in $Mflop/s$) of a particular block size on a sparse matrix A . We refer to the number of nonzeros in A as $k(A)$. We refer to the number of blocks of size $r \times c$ which contain nonzeros in A as $k_{r,c}(A)$. We can then define the *fill* of the matrix to be $f(A) = rck_{r,c}(A)/k(A)$. Once per machine, we compute a profile of how the machine performs for a particular block size. Let $P_{rc}(dense)$ be the performance of the machine (in $Mflop/s$) on a dense matrix stored with block size $r \times c$. Then we can estimate $P_{rc}(A)$ as

$$\tilde{P}_{rc}(A) = \frac{P_{rc}(dense)}{f_{rc}(A)}$$

Thus, our task is to compute f_{rc} for all r and c to within some tolerable relative accuracy, and to do so efficiently. Statistical sampling methods given by Vuduc et. al. provide no theoretical guarantee of accuracy, and take as long as 1 to 10 times the time it takes to perform a sparse matrix vector multiplication on the same matrix. We describe an algorithm which provides estimates to within ϵ relative error with probability $1 - \delta$ in time $O(\log(\delta)/\epsilon^2)$, and show that our algorithm runs efficiently and accurately on real-world cases. Note that our algorithm depends only on the desired accuracy, whereas the algorithm in [?] depends linearly upon the number of nonzeros.

Our algorithm estimates a more general notion of fill for tensors, where we divide the tensor into smaller subarrays (our blocks) and again only the nonzero blocks and their locations are stored in each row of the tensor. We further generalize this by allowing the user to offset the grid of blocks by some fixed amount, so that the block structure of the tensor does not have to align with the block size.

Finally, we note that estimating the fill can be an important part of any sparse datastructure which uses blocking, not just BCSR. In fact, any sparse datastructure can be adapted to a blocked regime by grouping a tensor into blocks and simply treating nonzero blocks as nonzeros of some sparse tensor.

1 Notation

A *tensor* is a multidimensional array. A tensor of *order* N is an element of the tensor (direct) product of N vector spaces. We assume all of our vector spaces are over an arbitrary field \mathbb{F} . Vectors are order 1 tensors and will be denoted by boldface lowercase letters, like this: \mathbf{a} . Matrices are order 2 tensors and will be denoted by boldface capital letters, like this: \mathbf{A} . Tensors will be denoted by boldface capital Euler script letters, like this: \mathcal{A} .

We will refer to the number of nonzero elements in a tensor \mathcal{A} by $k(\mathcal{A})$.

The n^{th} element in a sequence is denoted by $\mathcal{A}^{(n)}$. Element (i_1, \dots, i_N) of an order- N tensor $\mathcal{A} \in \mathbb{F}^{I_1 \times \dots \times I_N}$ is denoted $\mathcal{A}_{i_1, \dots, i_N}$ or $\mathcal{A}[i_1, \dots, i_N]$.

If we wish to represent the sequence $a, a + 1, \dots, b$, we use the syntax $a \rightarrow b$. When this sequence appears as a subscript, it is used as a shorthand for a sequence of subscripted

variables. For example, we use the syntax $i_{1 \rightarrow N}$ as a shorthand for i_1, \dots, i_N .

Subarrays are formed when we fix a subset of indices. We use a colon to indicate all elements of a mode. Thus, the middle $n/2$ columns of a matrix $\mathbf{A} \in \mathbb{F}^{n \times n}$ would be written $\mathbf{A}_{:,n/4 \rightarrow 3n/4}$.

For convenience, we say that $\mathcal{A} = 0$ if and only if every element of \mathcal{A} is 0.

We refer to populations using capital Euler script letters, like this: \mathcal{X} . We refer to random variables and index bounds using capital letters, like this: X . We refer to functions, indices and elements of populations using lowercase letters, like this x .

2 Formulation of the problem

We are given a tensor $\mathcal{A} \in \mathbb{F}^{I_1 \times \dots \times I_N}$ and a positive integer $B > 0$. We define the function $r_{b,o}$ for $1 \leq b, o \leq B$ as follows:

$$r_{b,o}(j) = (o + j * (b - 1) + 1) \rightarrow (o + j * b)$$

Intuitively, if we grouped the natural numbers into blocks of length b and then offset these blocks by o , $r_{b,o}(j)$ represents the range of indices corresponding to the j^{th} block.

In an N -dimensional setting,

$$r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N}) = r_{b_1, o_1}(j_1) \times \dots \times r_{b_N, o_N}(j_N)$$

We also define a corresponding function l that looks up the corresponding block index j of an index i .

$$l_{b,o}(i) = \left\lceil \frac{i}{b} \right\rceil$$

In an N -dimensional setting,

$$l_{b_1 \rightarrow N, o_1 \rightarrow N}(i_{1 \rightarrow N}) = (l_{b_1, o_1}(i_1), \dots, l_{b_N, o_N}(i_N))$$

We define $k_{b_1 \rightarrow N, o_1 \rightarrow N}$ for $1 \leq b_1, \dots, b_N, o_1, \dots, o_N \leq B$ as follows.

$$k_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}) = |\{(j_{1 \rightarrow N}) | \mathcal{A}[r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N})] \neq 0\}|$$

Thus, if we broke up our range of tensor indices into blocks of size b_1, \dots, b_N and offset these blocks by o_1, \dots, o_N , $k_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A})$ tells us how many of these blocks would be needed to cover the nonzeros of \mathcal{A} . Note that $k_{1, \dots, 1, 0, \dots, 0}(\mathcal{A}) = k(\mathcal{A})$.

Now we can formally define the fill $f_{b_1 \rightarrow N, o_1 \rightarrow N}$ for all $1 \leq b_1, \dots, b_N, o_1, \dots, o_N \leq B$.

$$f_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}) = \frac{k_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A})}{k(\mathcal{A})}$$

The problem is to compute approximate $\tilde{f}_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A})$ such that $f_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A})(1 - \epsilon) \leq \tilde{f}_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}) \leq f_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A})(1 + \epsilon)$ for all $1 \leq b_1, \dots, b_N, o_1, \dots, o_N \leq B$ with probability at least $1 - \delta$.

3 Previous Work

finish plz. Mainly this is Vuduc

4 The Algorithm

We define the function $x_{b_1 \rightarrow N, o_1 \rightarrow N}$ on each nonzero element $(i_{1 \rightarrow N})$ of \mathcal{A} as follows.

$$x_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}, i_{1 \rightarrow N}) = 1/k(\mathcal{A}[r_{b_1 \rightarrow N, o_1 \rightarrow N}(l_{b_1 \rightarrow N, o_1 \rightarrow N}(i_{1 \rightarrow N}))])$$

$x_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}, i_{1 \rightarrow N})$ is therefore equal to the reciprocal of the number of nonzeros in its block. Consider the sum of $x_{b_1 \rightarrow N, o_1 \rightarrow N}$ over all of the nonzeros of A . We have that

$$\begin{aligned} & \sum_{(i_{1 \rightarrow N}) | \mathcal{A}[i_{1 \rightarrow N}] \neq 0} x_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}, i_{1 \rightarrow N}) \\ &= \sum_{(j_{1 \rightarrow N}) | \mathcal{A}[r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N})] \neq 0} \left(\sum_{(i_{1 \rightarrow N}) \in r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N}) | \mathcal{A}[i_{1 \rightarrow N}] \neq 0} x_{b_1 \rightarrow N, o_1 \rightarrow N}(\mathcal{A}, i_{1 \rightarrow N}) \right) \\ &= \sum_{(j_{1 \rightarrow N}) | \mathcal{A}[r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N})] \neq 0} \sum_{(i_{1 \rightarrow N}) \in r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N}) | \mathcal{A}[i_{1 \rightarrow N}] \neq 0} 1/k(\mathcal{A}[r_{b_1 \rightarrow N, o_1 \rightarrow N}(j_{1 \rightarrow N})]) \\ &= \sum_{(j_{1 \rightarrow N}) | \mathcal{A}_{r_{b_1, o_1}(j_1), \dots, r_{b_N, o_N}(j_N)} \neq 0} 1 \\ &= \left| \{(j_{1 \rightarrow N}) | \mathcal{A}_{r_{b_1, o_1}(j_1), \dots, r_{b_N, o_N}(j_N)} \neq 0\} \right| \\ &= k_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A}) \end{aligned}$$

Consider the population $\mathcal{X}_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A}) = (x_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A}, i_{1 \rightarrow N}) | \mathcal{A}(i_{1 \rightarrow N}) \neq 0)$. Note that the average value of elements in $\mathcal{X}_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A})$ is

$$\frac{\sum_{(i_{1 \rightarrow N}) | \mathcal{A}(i_{1 \rightarrow N}) \neq 0} x_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A}, i_{1 \rightarrow N})}{\|\{(i_{1 \rightarrow N}) | \mathcal{A}(i_{1 \rightarrow N}) \neq 0\}\|} = \frac{k_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A})}{k(\mathcal{A})} = f_{b_1, \dots, b_N, o_1, \dots, o_N}(\mathcal{A})$$

Thus, our task is to randomly sample elements from $\mathcal{X}_{b_1, \dots, b_N, o_1, \dots, o_N}$ to compute an estimate of its average. We can compute a sample of $\mathcal{X}_{b_1, \dots, b_N, o_1, \dots, o_N}$ by randomly choosing counting the number of nonzeros Randomly sampling a particular element of $XX_{b_1, \dots, b_N, o_1, \dots, o_N}$ can be difficult depending on how the sparse array is stored. However, we can demonstrate a method for sampling $q_{b_1, \dots, b_N, o_1, \dots, o_N}$

5 Analysis of Algorithm

Here, we will beat the analysis of this algorithm to death so that we can get a bound on the number of operations required to compute this estimate (we are talking about constants after all)

5.1 Error Analysis

Here, we will bound (very tightly) the number of samples needed

5.2 Runtime Analysis

Here, we will bound (very tightly) the number of operations per sample needed

6 Results

Here we explore the relationship between runtime and accuracy of the fill prediction on several matrices from Vuduc et. al. and also from the suitesparse collection from florida